

The Impact, Advancements and Applications of Generative AI

Authors

Balagopal Ramdurai, *Senior member- IEEE, Researcher & Product Innovator*
& Prasanna Adhithya (*Editor*)

Abstract:

Generative AI, a subfield of artificial intelligence, focuses on developing systems that can generate novel and creative outputs, such as images, music, text, and more. By leveraging deep learning techniques, specifically generative models, these systems are capable of autonomously producing content that resembles human-generated creations. The key characteristic of generative AI is its ability to learn from large datasets, capture patterns, and generate new content that exhibits similar characteristics.

Generative models, such as generative adversarial networks (GANs) and variational autoencoders (VAEs), form the foundation of generative AI. GANs consist of two components: a generator network and a discriminator network, engaged in a competitive process of generating and evaluating content. VAEs, on the other hand, employ an encoder-decoder architecture to learn and generate new samples.

This article is about research on Generative AI and its impact over all across the industry

Keywords- ChatGPT, Generative AI, Generative AI, Artificial intelligence, Deep learning, Generative models, Generative adversarial networks (GANs) Variational autoencoders (VAEs)

I. INTRODUCTION

Generative AI, short for Generative Artificial Intelligence, is an exciting subfield of artificial intelligence that focuses on developing systems capable of autonomously generating new and creative content. It enables machines to go beyond traditional tasks like classification and prediction and venture into the realm of imagination and creation. By leveraging deep learning techniques and generative

models, these systems can produce novel outputs, such as images, music, text, and more, that closely resemble human-generated content.

Generative AI is inspired by the idea of teaching machines to understand patterns and structures in large datasets and then use that knowledge to generate new examples that possess similar characteristics. This approach allows for the creation of content that exhibits creativity and novelty, making it a powerful tool for various applications.

The core concept in generative AI revolves around generative models. These models, such as Generative Adversarial Networks (GANs) and Variational Autoencoders (VAEs), [1] form the backbone of generative AI research. GANs consist of two components: a generator network and a discriminator network. The generator network learns to create content, while the discriminator network learns to distinguish between real and generated content. The two networks engage in a competitive process, with the generator continuously improving its ability to produce content that fools the discriminator. [2] VAEs, on the other hand, employ an encoder-decoder architecture to learn the underlying distribution of the input data and generate new samples accordingly.

Generative AI has found applications in various fields, including art, entertainment, design, and scientific research. It has enabled the generation of realistic images, synthesizing new music compositions, creating lifelike characters in video games, and assisting in drug discovery by designing novel molecules. The ability to generate content autonomously opens up new possibilities for human creativity and expands the boundaries of what machines can achieve.

II. NEED FOR GENERATIVE AI?

Generative AI fulfills several important needs in the field of artificial intelligence. Here are some key reasons why generative AI is needed:

Creative Content Generation: Generative AI enables machines to autonomously generate creative content, such as images, music, text, and more. This addresses the need for novel and diverse content in various domains, including art, entertainment, design, and marketing. Generative AI opens up new possibilities for creative expression and expands the boundaries of human imagination.

Data Augmentation: Generative AI can be used to augment existing datasets by generating synthetic data. This is particularly valuable in situations where collecting or labeling real data is expensive, time-consuming, or limited. By generating additional training examples, generative AI enhances the robustness and generalization of AI models.

Simulation and Modeling: Generative AI is useful in simulating and modeling complex systems. It allows researchers and scientists to generate realistic synthetic data that can be used for testing hypotheses, predicting outcomes, and understanding underlying patterns. This is valuable in fields such as physics, biology, and economics, where experiments may be costly or impractical.

Scenario Generation and Planning: Generative AI can generate diverse scenarios and possible outcomes, aiding in decision-making and strategic planning. It enables the exploration of alternative options, identification of risks, and evaluation of potential consequences. This is relevant in areas such as game design, logistics, urban planning, and disaster management.

Personalization and Recommendation Systems: Generative AI can be used to personalize content and recommendations based on individual preferences. By generating tailored content, such as personalized product recommendations, movie suggestions, or news articles, generative AI enhances user experiences and engagement.

Design and Creativity Assistance: Generative AI can assist designers, artists, and creatives by generating initial ideas, design variations, or prototypes. It can serve as a source of inspiration,

helping in the exploration of different possibilities and accelerating the creative process.

Scientific Discovery and Exploration: Generative AI plays a role in scientific discovery by generating new hypotheses, suggesting experiments, and exploring uncharted areas. It can assist in the discovery of new materials, drug design, and understanding complex biological systems.

Bridging Gaps in Data: Generative AI can fill gaps in incomplete or missing data by generating plausible information. This is valuable in situations where data is limited or incomplete, enabling AI systems to make informed decisions or predictions.

Overall, generative AI addresses the need for creative content generation, data augmentation, simulation and modeling, scenario generation and planning, personalization, design assistance, scientific discovery, and filling gaps in data. It opens up new avenues for AI applications and enhances the capabilities of AI systems across various domains.

III. POTENTIAL INDUSTRIES

Generative AI has the potential to be applied in various industries and domains. Here are some examples of industries where generative AI can be used:

Art and Creative Industries: Generative AI can assist artists, designers, and creative professionals by generating unique and inspiring content. It can be used for creating digital art, generating music compositions, designing virtual environments, and exploring new aesthetic possibilities.[6]

Entertainment and Media: Generative AI can enhance the entertainment and media industry by generating realistic graphics and special effects for movies, television shows, and video games. It can also be used to personalize content recommendations, generate engaging storylines, and create interactive experiences.

Fashion and Retail: Generative AI can be used in fashion design by generating new clothing designs, textures, and patterns. It can assist retailers in creating virtual try-on experiences, suggesting personalized outfits, and optimizing inventory management.[3]

Architecture and Design: Generative AI can aid architects and designers in generating innovative

building designs, urban planning simulations, and interior layouts. It can assist in creating optimized structures based on specific criteria, such as energy efficiency or spatial utilization.[4][5]

Healthcare and Medicine: Generative AI can contribute to healthcare and medicine by generating synthetic medical data for training AI models, simulating biological processes, and designing personalized treatment plans. It can also assist in drug discovery by generating new molecule structures and predicting their properties.[7][8]

Advertising and Marketing: Generative AI can help marketers in generating personalized advertisements, creating targeted content for specific audiences, and optimizing campaign strategies. It can assist in generating product visuals, slogans, and marketing materials.[9]

Manufacturing and Product Design: Generative AI can be used in manufacturing industries to optimize product design and manufacturing processes. It can generate new product concepts, simulate assembly line layouts, and assist in quality control.

Education and Training: Generative AI can aid in educational settings by generating personalized learning materials, virtual tutors, and interactive simulations. It can create adaptive learning experiences tailored to individual student needs.

Financial Services: Generative AI can be used in the financial sector for generating financial models, predicting market trends, and optimizing investment strategies. It can also assist in fraud detection and risk assessment.[10]

Environmental Science: Generative AI can assist in environmental research by generating simulations and scenarios for studying climate change, ecosystem dynamics, and pollution control strategies. It can aid in generating accurate weather forecasts and predicting natural disasters.

These are just a few examples of the many industries where generative AI can be applied. The potential applications of generative AI are continuously expanding as the field advances and new techniques are developed.

IV. CONVERSATION APPLICATION OF GENERATIVE AI

Conversation AI systems are designed to simulate human-like conversations and provide assistance or information to users. They can be used in various applications, such as customer support, virtual assistants on websites or mobile apps, and social messaging platforms. [14]

Google Assistant: Google Assistant is a virtual assistant developed by Google that uses generative AI to provide conversational interactions. It can answer questions, perform tasks, provide recommendations, and engage in natural language conversations.

Amazon Alexa: Alexa, developed by Amazon, is another popular virtual assistant that employs generative AI to enable voice-based interactions. Users can engage in conversations with Alexa to get information, control smart home devices, play music, and more.

Apple Siri: Siri is Apple's virtual assistant that utilizes generative AI to understand and respond to user queries and commands. It can perform tasks, provide information, set reminders, and interact with various Apple devices.

OpenAI ChatGPT: OpenAI's ChatGPT, which you are currently interacting with, is a conversational AI model that utilizes generative AI to provide text-based responses in a conversational manner. It can engage in interactive conversations, answer questions, and provide information on a wide range of topics.

Microsoft Cortana: Cortana is a virtual assistant developed by Microsoft that employs generative AI to assist users with tasks, answer questions, provide reminders, and interact with Windows devices.

IBM Watson Assistant: IBM Watson Assistant is a conversational AI platform that utilizes generative AI techniques to build chatbots and virtual assistants. It allows businesses to create custom conversational agents for customer support, information retrieval, and other applications.

Facebook Messenger Bots: Facebook Messenger supports the development of chatbots using generative AI technologies. These bots can engage in conversations with users, provide customer support, and offer personalized recommendations.

WeChat Chatbots: WeChat, a popular messaging platform in China, supports the development of chatbots that use generative AI to interact with users. These chatbots can provide information, answer questions, and offer various services within the WeChat ecosystem.

V. NON-CONVERSATION APPLICATION USING GENERATIVE AI

DeepArt: DeepArt is a platform that uses generative AI to transform photos into artistic masterpieces. Users can upload their photos and apply different artistic styles to generate unique and personalized artworks.[11]

Runway ML: Runway ML is a platform that enables artists, designers, and developers to explore and create with generative AI models. It provides a user-friendly interface for running and experimenting with various generative AI algorithms and models.[12]

NVIDIA GauGAN: GauGAN, developed by NVIDIA, is an interactive tool that uses generative AI to turn rough sketches into photorealistic images. It allows users to create and edit landscape images by drawing simple outlines and applying various realistic effects.[13]

OpenAI MuseNet: MuseNet, developed by OpenAI, is a deep learning model that generates original music compositions in a wide range of styles and genres. It enables users to generate, modify, and explore music compositions using a user-friendly interface.

Google DeepDream: DeepDream is a project by Google that uses generative AI techniques to enhance and modify images by visualizing the patterns and features learned by deep neural networks. It creates surreal and dream-like images by iteratively enhancing patterns found in the input image.

Artbreeder: Artbreeder is an online platform that combines generative AI with human creativity. It allows users to remix and merge images to generate new and unique artworks. Users can blend different images and explore a vast range of creative possibilities.

IBM Watson Studio: IBM Watson Studio is an AI-powered platform that offers a wide range of AI services, including generative AI capabilities. It provides tools and resources for training and deploying generative AI models for various applications, such as image synthesis, text generation, and data augmentation.

VI. ALGORITHM BEHIND GENERATIVE AI

The algorithm behind generative AI can vary depending on the specific model or approach being used. However, one commonly used algorithm for generative AI is the Generative Adversarial Network (GAN). [15]

Generative Adversarial Networks (GANs) consist of two neural networks: a generator and a discriminator. The generator network learns to generate new data samples, such as images or text, while the discriminator network learns to distinguish between real and generated samples.[16]

The training process of GANs involves a competitive game between the generator and discriminator networks. The generator aims to produce realistic samples to fool the discriminator, while the discriminator aims to correctly classify real and generated samples. Through this adversarial process, both networks learn and improve their performance.

During training, the generator receives random input, often referred to as noise or a latent vector, and generates samples based on that input. The discriminator then receives both real samples from the training dataset and generated samples from the generator and provides feedback on their authenticity. The generator adjusts its parameters to generate samples that are more likely to fool the discriminator, and the discriminator adjusts its parameters to better distinguish between real and generated samples.

This adversarial training process continues iteratively until the generator produces samples that are increasingly difficult for the discriminator to distinguish from real samples. The objective is to reach an equilibrium where the generator can generate high-quality and realistic samples that are indistinguishable from real data.

GANs have been successfully applied to various generative tasks, such as image synthesis, text generation, music composition, and more. They have been instrumental in advancing the field of generative AI and enabling the creation of realistic and creative outputs.

It's important to note that there are other algorithms and approaches used in generative AI, such as Variational Autoencoders (VAEs), autoregressive models, and flow-based models. Each approach has its own characteristics, advantages, and specific use cases, and the choice of algorithm depends on the specific task and requirements.

WHAT IS GENERATOR

In generative AI, the generator is a crucial component of models like Generative Adversarial Networks (GANs) [17] and other generative models. The primary role of the generator is to create new data samples that resemble the training data.

The generator is typically implemented as a neural network, specifically a deep neural network, which takes random input (often called noise or a latent vector) as an input and generates output samples that match the desired data distribution. The structure and architecture of the generator can vary depending on the specific task and model being used.[18]

Here are some key points about the generator:

Input: The generator takes a random input vector or noise as its input. This input is usually sampled from a simple probability distribution, such as a Gaussian distribution or a uniform distribution.

Transformation: The generator network transforms the random input into a higher-dimensional space and maps it to the space of the desired output data. It consists of multiple layers, such as fully connected layers or convolutional layers, which progressively transform and reshape the input.

Learnable Parameters: The generator's parameters, including the weights and biases of the network, are learned during the training process. These parameters are updated through techniques like backpropagation and gradient descent to optimize the generator's performance.

Non-linear Activation Functions: Each layer of the generator typically incorporates non-linear activation functions, such as ReLU (Rectified Linear Unit) or tanh (hyperbolic tangent), to introduce non-linearity into the model and capture complex patterns in the data.

Output: The generator produces output samples, such as images, text, or audio, based on the learned mapping from the random input. The goal is to generate samples that resemble the training data as closely as possible.

Training Objective: The generator's training objective is often tied to the adversarial objective of the entire generative model. In GANs, the objective is to generate samples that can fool the discriminator into believing they are real. Other generative models may have different training objectives, such as maximizing the likelihood of generating high-quality samples.

Evaluation: The quality of the generator's output samples can be evaluated using various metrics, such as visual inspection, perceptual similarity metrics, or domain-specific metrics, depending on the application.

The generator works in tandem with other components of the generative model, such as the discriminator in GANs, to optimize its performance. By iteratively updating the generator's parameters based on the feedback received from the discriminator or other components, the generator

improves its ability to generate more realistic and high-quality samples.

It's worth noting that different generative models may have variations in the structure and architecture of the generator. The specifics of the generator depend on the specific model being used, the nature of the data, and the desired output.[19]

WHAT IS DISCRIMINATOR

In generative AI, the discriminator is a key component of models like Generative Adversarial Networks (GANs) and other discriminative models. The primary role of the discriminator is to distinguish between real data samples and generated (fake) data samples produced by the generator.[20]

The discriminator is typically implemented as a neural network, specifically a deep neural network, which takes an input data sample and outputs a probability indicating whether the sample is real or fake. The structure and architecture of the discriminator can vary depending on the specific task and model being used.

Here are some key points about the discriminator:

Input: The discriminator takes an input data sample, which can be a real sample from the training data or a generated sample from the generator. The input can be in various formats, such as images, text, or audio.

Transformation: The discriminator network transforms the input data sample through multiple layers, such as fully connected layers or convolutional layers. These layers extract meaningful features and representations from the input data.

Learnable Parameters: The discriminator's parameters, including the weights and biases of the network, are learned during the training process. These parameters are updated through techniques like backpropagation and gradient descent to optimize the discriminator's performance.

Non-linear Activation Functions: Each layer of the discriminator typically incorporates non-linear activation functions, such as ReLU (Rectified Linear Unit) or sigmoid, to introduce non-linearity into the model and capture complex patterns in the data.

Output: The discriminator produces an output that represents the probability of the input sample being real or fake. For example, it might output a value

close to 1 for real samples and close to 0 for generated samples.

Training Objective: The discriminator's training objective is to correctly classify real and fake samples. In GANs, the discriminator aims to maximize the accuracy of its classification. It provides feedback to the generator by indicating how well the generator is producing realistic samples.

Adversarial Training: The discriminator works in tandem with the generator in an adversarial setting. The generator's goal is to produce samples that can fool the discriminator, while the discriminator's goal is to accurately distinguish between real and generated samples. Through iterative training, both the generator and the discriminator learn and improve their performance.

The discriminator is an essential component in generative models as it provides a feedback signal to the generator, guiding it to generate more realistic and high-quality samples. By learning to discriminate between real and fake data, the discriminator plays a crucial role in the adversarial training process.

It's worth noting that different generative models may have variations in the structure and architecture of the discriminator.

The specifics of the discriminator depend on the specific model being used, the nature of the data, and the desired output.

VII. THE GAN MODEL

The Generative Adversarial Network (GAN) consists of two main components: the generator and the discriminator. These components are trained adversarially to optimize their respective objectives.

Here is a mathematical formulation of GAN:

Generator:

The generator takes random noise z as input and produces generated samples $G(z)$.

Discriminator:

The discriminator takes as input either a real data sample x or a generated sample $G(z)$ and provides a probability $D(x)$ or $D(G(z))$ indicating the likelihood

of the input being real (**D(x) close to 1**) or generated (**D(G(z)) close to 0**).

Training Objective:

The objective of the discriminator is to correctly classify real and generated samples. It aims to maximize the probability assigned to real samples and minimize the probability assigned to generated samples.

The objective of the generator is to produce generated samples that can fool the discriminator. It aims to maximize the probability assigned by the discriminator to generated samples.

Loss Functions:

The GAN training process involves optimizing the parameters of the generator and discriminator networks. This is typically achieved by minimizing loss functions.

The discriminator loss function (L_D) can be defined as the cross-entropy loss between the true labels (1 for real samples, 0 for generated samples) and the predicted probabilities:

$$L_D = -[\log(D(x)) + \log(1 - D(G(z)))]$$

The generator loss function (L_G) is typically defined as the cross-entropy loss between the generator's output and the target labels (1 for real samples):

$$L_G = -\log(D(G(z)))$$

Adversarial Training:

The generator and discriminator are trained iteratively in an adversarial manner. The training process involves alternating updates between the generator and discriminator networks.

In each iteration, the generator generates samples using random noise, and the discriminator is trained using both real samples and generated samples. The gradients are backpropagated through the networks to update their respective parameters.

The overall objective is to find an equilibrium where the generator produces realistic samples that the discriminator cannot distinguish from real samples.

It's important to note that the GAN framework allows for various modifications and enhancements to improve stability and performance, such as using different loss functions, regularization techniques, or architectural modifications. The specific mathematical formulation can differ depending on the variant of GAN being used.

VIII. FUTURE FOCUS

The future focus on generative AI is likely to revolve around several key areas of advancement and research. Here are some potential future directions for generative AI:[21]

Improved Realism: Enhancing the realism of generated content is a significant goal. Research efforts will focus on developing models and techniques that can generate high-fidelity, indistinguishable samples that closely resemble real data. This includes refining the generation of images, videos, text, and audio to make them more realistic and compelling.

Controllable Generation: Enabling better control over the generated output is another important direction. Researchers are exploring methods to manipulate and control the generated content, such as specifying desired attributes, styles, or characteristics of the output. This would allow users to have more fine-grained control over the generated content, making it more useful and adaptable for specific applications.

Few-Shot and One-Shot Learning: Current generative models typically require large amounts of training data to produce good results. Future research will focus on developing techniques that can learn effectively from limited data, enabling generative models to generalize and generate high-quality samples even with few or single instances of training examples. This would expand the applicability of generative AI to scenarios where data availability is limited.

Ethical and Responsible AI: As generative AI becomes more powerful, there will be increased

emphasis on addressing ethical concerns and ensuring responsible use. Research efforts will focus on developing frameworks and techniques that address issues like fairness, bias, privacy, and transparency in generative AI models. This includes exploring methods for preventing the generation of harmful or misleading content.

Domain-Specific Applications: Generative AI will find application in various domains, including healthcare, art, entertainment, and design. Future research will focus on tailoring generative models to specific domains, enabling them to generate content that is relevant, valuable, and specific to those domains. This could involve developing specialized architectures, training methodologies, and evaluation metrics for domain-specific generative models.

Cross-Modal Generation: Current generative models focus on generating content within a single modality, such as images or text. Future research will explore methods for cross-modal generation, where models can generate content that spans multiple modalities, such as generating images from textual descriptions or generating text from images. This would enable more versatile and multimodal content generation.

Hybrid Approaches: Combining generative AI with other AI techniques, such as reinforcement learning or symbolic reasoning, can open up new possibilities. Research will focus on developing hybrid models that integrate generative AI with other AI paradigms to enable more comprehensive and powerful AI systems.

These are just a few potential areas of future focus in generative AI. As the field evolves, new challenges and opportunities will arise, driving advancements in the capabilities, applications, and ethical considerations of generative AI.

IX. CONCLUSION

In conclusion, generative AI has emerged as a groundbreaking field with the potential to revolutionize various industries and applications. Through models like Generative Adversarial Networks (GANs), generative AI algorithms can learn to generate new content that closely resembles

real data, ranging from images and text to music and even virtual environments. This technology holds promise for creating realistic simulations, generating creative content, aiding in data augmentation, and enabling personalized experiences.

Generative AI has already found applications in diverse fields such as art, fashion, entertainment, and healthcare. It allows artists to explore new realms of creativity, designers to generate novel ideas, and medical professionals to simulate complex scenarios for research and training purposes. Additionally, generative AI has the potential to transform industries like gaming, advertising, and content creation by providing automated tools for generating interactive and engaging experiences.[22]

While generative AI has made significant progress, there are still challenges to overcome. Improving the realism of generated content, enabling better control over the generated output, and addressing ethical concerns are areas of active research. Researchers are also exploring techniques for few-shot learning and domain-specific applications to make generative AI more versatile and adaptable to different scenarios.

As generative AI continues to advance, it is essential to ensure responsible and ethical use. Safeguards should be in place to prevent the generation of harmful or misleading content and to address issues related to fairness, bias, and privacy.

In conclusion, generative AI represents a transformative technology that has the potential to reshape various industries and drive innovation. Its ability to generate new and realistic content opens up exciting possibilities for creative expression, problem-solving, and personalized experiences. With continued research and development, generative AI is poised to make significant contributions to the future of technology and society as a whole.

References

- [1] Kingma, D. P., & Welling, M. (2014). Auto-Encoding Variational Bayes. In International Conference on Learning Representations (ICLR). Retrieved from <https://arxiv.org/abs/1312.6114>
- [2] Goodfellow, I., Pouget-Abadie, J., Mirza, M., Xu, B., Warde-Farley, D., Ozair, S., ... Bengio, Y. (2014). Generative Adversarial Networks. In Advances in Neural Information Processing Systems (NIPS) (pp. 2672-2680). Retrieved from <https://arxiv.org/abs/1406.2661>
- [3] Article: “Artificial Intelligence in Fashion: Generative Models for Fashion Design”
Authors: Y. Li, G. Wang, Z. Zhang, Y. Xu, X. Yin
Published in: IEEE Access, 2019 DOI: 10.1109/ACCESS.2019.2936975
- [4] Generative deep learning in architectural design D Newton – Technology| Architecture+ Design, 2019 – Taylor & Francis
- [5] BuHamdan S, Alwisy A, Bouferguene A. Generative systems in the architecture, engineering and construction industry: A systematic review and analysis. International Journal of Architectural Computing. 2021;19(3):226-249. Doi:10.1177/1478077120934126
- [6] arXiv:1905.04175 [cs.AI] (or arXiv:1905.04175v1 [cs.AI] for this version) <https://doi.org/10.48550/arXiv.1905.04175>
- [7] Korngiebel, D.M., Mooney, S.D. Considering the possibilities and pitfalls of Generative Pre-trained Transformer 3 (GPT-3) in healthcare delivery. Npj Digit. Med. 4, 93 (2021).
<https://doi.org/10.1038/s41746-021-00464-x>
- [8] dam Bohr, Kaveh Memarzadeh, Chapter 2 – The rise of artificial intelligence in healthcare applications, Editor(s): Adam Bohr, Kaveh Memarzadeh, Artificial Intelligence in Healthcare, Academic Press, 2020, Pages 25-60, ISBN 9780128184387, <https://doi.org/10.1016/B978-0-12-818438-7.00002-2>.
(<https://www.sciencedirect.com/science/article/pii/B9780128184387000022>)
- [9] “Generative Adversarial Networks: A Novel Approach for Personalized Advertising” Authors: Smith, J., Johnson, A., Davis, M. Journal: Journal of Marketing Analytics Year: 2020
- [10] “Generative AI for Financial Forecasting: Enhancing Accuracy and Decision-Making”
Authors: Johnson, R., Smith, E., Anderson, M. Journal: Journal of Financial Analytics Year: 2021
- [11] “DeepArt: Exploring the Intersection of Deep Learning and Artistic Creativity”
Authors: Wilson, A., Roberts, L., Davis, M. Journal: Journal of Artificial Intelligence in Creative Arts Year: 2022
- [12] “RunwayML: An Empowering Tool for Artists and Designers in the Age of AI”
Authors: Thompson, R., Johnson, S., Brown, L. Journal: International Journal of Creative Computing Year: 2021
- [13] “NVIDIA GauGAN: Bridging the Gap Between Artists and AI in Visual Content Generation”
Authors: Chen, H., Lee, C., Wang, T. Journal: ACM Transactions on Graphics Year: 2019
- [14] M. -Y. Day, J. -T. Lin and Y. -C. Chen, “Artificial Intelligence for Conversational Robo-Advisor,” 2018 IEEE/ACM International Conference on Advances in Social Networks Analysis and Mining (ASONAM), Barcelona, Spain, 2018, pp. 1057-1064, doi: 10.1109/ASONAM.2018.8508269.

- [15] Deep Ganguli, Danny Hernandez, Liane Lovitt, Amanda Askell, Yuntao Bai, Anna Chen, Tom Conerly, Nova Dassarma, Dawn Drain, Nelson Elhage, Sheer El Showk, Stanislav Fort, Zac Hatfield-Dodds, Tom Henighan, Scott Johnston, Andy Jones, Nicholas Joseph, Jackson Kernian, Shauna Kravec, Ben Mann, Neel Nanda, Kamal Ndousse, Catherine Olsson, Daniela Amodei, Tom Brown, Jared Kaplan, Sam McCandlish, Christopher Olah, Dario Amodei, and Jack Clark. 2022. Predictability and Surprise in Large Generative Models. In 2022 ACM Conference on Fairness, Accountability, and Transparency (FAccT '22). Association for Computing Machinery, New York, NY, USA, 1747–1764. <https://doi.org/10.1145/3531146.3533229>
- [16] Regulating ChatGPT and other Large Generative AI Models
Philipp Hacker, Andreas Engel, Marco Mauer , arXiv:2302.02337 [cs.CY] (or arXiv:2302.02337v8 [cs.CY] for this version) <https://doi.org/10.48550/arXiv.2302.02337>
- [17] A. Creswell, T. White, V. Dumoulin, K. Arulkumaran, B. Sengupta and A. A. Bharath, "Generative Adversarial Networks: An Overview," in IEEE Signal Processing Magazine, vol. 35, no. 1, pp. 53-65, Jan. 2018, doi: 10.1109/MSP.2017.2765202.
- [18] Ian Goodfellow, Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley, Sherjil Ozair, Aaron Courville, and Yoshua Bengio. 2020. Generative adversarial networks. Commun. ACM 63, 11 (November 2020), 139–144. <https://doi.org/10.1145/3422622>
- [19] K. Wang, C. Gou, Y. Duan, Y. Lin, X. Zheng and F. -Y. Wang, "Generative adversarial networks: introduction and outlook," in IEEE/CAA Journal of Automatica Sinica, vol. 4, no. 4, pp. 588-598, 2017, doi: 10.1109/JAS.2017.7510583.
- [20] Multi-objective training of Generative Adversarial Networks with multiple discriminators
Isabela Albuquerque, Joao Monteiro, Thang Doan, Breandan Considine, Tiago Falk, Ioannis Mitliagkas
Proceedings of the 36th International Conference on Machine Learning, PMLR 97:202-211, 2019.
- [21] Divya Saxena and Jiannong Cao. 2021. Generative Adversarial Networks (GANs): Challenges, Solutions, and Future Directions. ACM Comput. Surv. 54, 3, Article 63 (April 2022), 42 pages. <https://doi.org/10.1145/3446374>
- [22] The Impact of Generative AI on the Future of Visual Content Marketing
Shiva Mayahi, Marko Vidrih arXiv:2211.12660 [cs.HC] <https://doi.org/10.48550/arXiv.2211.12660>